

# Linear Regression with Multiple Variables: A Machine Learning Approach

Renzo A. Viloche Morales

June 25, 2017

## 1 General Model and Notation

Linear regression models are among the simplest prediction models for a single output based on multiple input variables or *features*<sup>1</sup>. It is assumed that the output is a *continuous* and a linear function of all considered inputs. Specifically, the general *hypothesis function* relates the predicted output  $\hat{y}$  and  $n$  features  $x_1, x_2, \dots, x_n$  with the following expression,

$$\hat{y} = h_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \dots + \theta_n x_n \quad (1)$$

As an alternative form, equation (1) can also be expressed as the inner product,

$$h_{\theta}(\mathbf{x}) = \boldsymbol{\theta}^T \cdot \mathbf{x} \quad (2)$$

where,  $\boldsymbol{\theta} = [\theta_0 \ \theta_1 \ \dots \ \theta_n]^T$  is a *parameter vector*, and  $\mathbf{x} = [x_0 \ x_1 \ \dots \ x_n]^T$  represent all the observed features<sup>2</sup>.

Both equations (1) and (2) show the solution of the *direct* or *forward problem*: finding out a prediction value from any set of observed features. Therefore, solving the forward problem implies that the relation between input and output can be “deterministically” defined. However, in practice there is little or no knowledge of the input-output mathematical description. One must first solve an *inverse problem* by finding out, for a given model, the parameters that best describe the available input and output observations [Keller, 1976, Kirsch, 2011]. In particular, the choice of a general model such as (1) leads to what is known as a *linear regression problem*.

As already mentioned, the central part of a linear regression relies on data availability. Consider a collection of  $m$  independent training examples, *e.g.*, a dataset comprehending a particular *target* (output observations) and the corresponding inputs. Assume now that the solution of the linear regression problem is already known (the set of parameters values that ‘best describe’ the input and target are known). This means that for every training input example, the predicted value by (1) is approximatively equal to the respective target value,

$$\mathbf{y} \approx \mathbf{X}\boldsymbol{\theta} \quad (3)$$

The expression (3) is an informal way of defining the solution characteristic of the linear regression problem, and therefore, does not represent a system of linear equations. Its right-hand side represents the predicted values using the vectorized form from the equation (2) when considering all  $m$  training examples. The left-hand side term represents the target in a vector form.

---

<sup>1</sup>This procedure is also known in the literature as *Multivariate Linear Regression*.

<sup>2</sup>In this notation, the first component satisfies  $x_0 = 1$  in order to be consistent with equation (1).

The notation considered in this work is that training inputs are stored at the  $m \times (n+1)$   $\mathbf{X}$  matrix and the target is represented by the  $\mathbf{y}$  column vector

$$\mathbf{X} = \begin{bmatrix} x_0^{(1)} & x_1^{(1)} & x_2^{(1)} & \dots & x_n^{(1)} \\ x_0^{(2)} & x_1^{(2)} & x_2^{(2)} & \dots & x_n^{(2)} \\ x_0^{(3)} & x_1^{(3)} & x_2^{(3)} & \dots & x_n^{(3)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_0^{(m)} & x_1^{(m)} & x_2^{(m)} & \dots & x_n^{(m)} \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ y^{(3)} \\ \vdots \\ y^{(m)} \end{bmatrix} \quad (4)$$

Accordingly, the  $i$ -th row vector  $\mathbf{x}^{(i)} = [x_0^{(i)} \ x_1^{(i)} \ \dots \ x_n^{(i)}]$  represents the  $i$ -th input sample, whereas the  $j$ -th column vector  $\mathbf{x}_j$  indicates all the collected samples for this specific feature.

The next section will focus on how to find the parameters in order that the “error” between predictions and target becomes acceptable as expressed by (3). In fact, this loose notion of error will be mathematically defined and used as a criterion for approaching the solution iteratively. In addition, an analytical method will be also presented and the benefits and drawbacks on performing a linear regression considering both approaches will be briefly discussed.

## 2 Minimizing the Cost Function

A linear regression can be equivalently stated as an optimization problem where a *cost function* or *objective function* must be minimized. The straightforward choice for the cost function is the quadratic sum of all residuals from output and the correspondent predicted value as the following expression,

$$J(\boldsymbol{\theta}) = \frac{1}{2m} \sum_{i=1}^m (h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)})^2 \quad (5)$$

where  $y^{(i)}$  is the  $i$ th target example and  $h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)})$  the  $i$ th predicted value. Equation (5) is just a general error measure that captures the overall deviation on the pair prediction-target<sup>3</sup>.

The optimal solution is represented by the set of parameters  $\theta_0, \theta_1, \theta_2, \dots, \theta_n$  that minimizes the expression (5). In formal terms, the optimization problem to be solved is therefore,

$$\min_{\boldsymbol{\theta} \in \mathbb{R}^{n+1}} \frac{1}{2m} \sum_{i=1}^m (h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)})^2 \quad (6)$$

A global minimum solution is always guaranteed to exist if the general model is used for computing predictions. A more general condition to the convergence of algorithms to the minimum is that  $J(\boldsymbol{\theta})$  must be a *convex function*. The next section will describe an iterative algorithm that can be used to find the parameter values.

---

<sup>3</sup>In fact, equation (5) is closely related to the mean square error (MSE). The factor of  $1/2$  is conveniently considered here as it simplifies the analytical expression on the Gradient Descent algorithm.

## 2.1 Gradient Descent Algorithm

The Gradient Descent algorithm operates iteratively computing at each step a solution that is (in theory) closer to the global minimum. The key idea here is to update the previous solution with the corresponding cost function opposite highest rate

$$\boldsymbol{\theta}_{new} = \boldsymbol{\theta}_{old} - \alpha \nabla[J(\boldsymbol{\theta}_{old})] \quad (7)$$

where the scalar  $\alpha$  is known as the *learning rate*. The reason for adopting a scalar (usually different from zero) is that the learning rate is able to module the displacement or “jumps” allowing a faster (or slower) convergence to the global minimum. Assuming the general linear regression model, and the cost function previously defined, the solution components are specifically updated via (7) as

$$\begin{aligned} \theta_0 &:= \theta_0 - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) \\ \theta_1 &:= \theta_1 - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})x_1^{(i)} \\ \theta_2 &:= \theta_2 - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})x_2^{(i)} \\ &\vdots \\ \theta_n &:= \theta_n - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})x_n^{(i)} \end{aligned} \quad (8)$$

where, the ‘:=’ symbol denotes the assignment operation commonly used for algorithms<sup>4</sup>. The Algorithm 1 describes the complete procedure of finding a solution to a Multiple Linear Regression<sup>5</sup>.

---

**Algorithm 1** Gradient Descent Algorithm (Multiple Linear Regression):

---

- 1: Set initial guess for the parameter values:  $\theta_0, \theta_1, \theta_2, \dots, \theta_n$
  - 2: Set the learning rate value  $\alpha$
  - 3: Set the total number of update iterations  $N_{iter}$
  - 4: **for**  $k = 1, 2, \dots, N_{iter}$  **do**
  - 5:   Compute the parameters (components must be simultaneously updated!),  
 $\theta_j := \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})x_j^{(i)}, \quad j = 0, 1, \dots, n$
  - 6:   Compute the cost,  
 $J(k) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})^2$
  - 7: **end for**
- 

Although not a required condition, a random set of values for the parameters can be used to initialize the search for the solution. The learning rate  $\alpha$  must be properly chosen in the sense that the algorithm produces a cost function that decreases with the number of iterations (the error is minimized). A rule of thumb says that the smaller the value of  $\alpha$  is, the greater the chance for the global minimum being reached (at the expense of a required higher number of iterations). In the other hand, a large learning rate can lead to convergence failure. In effect, large “jumps” for the solution update (equation 7) can produce a (i) divergent, or (ii) oscillating characteristic for the cost function evolution. Consequently, one must always check the  $J_k(\boldsymbol{\theta}) \times k$  behavior in order to assure an acceptable  $\boldsymbol{\theta}$  solution (this is why it is important to compute and store the error value at each iteration). In some cases, the number of iterations can be extended if a correct minimization procedure holds true, and a lower error value is desired.

---

<sup>4</sup>This means that the value on the right-side of this symbol will substitute any previous value present at the variable whose storing the  $\theta_i$  information.

<sup>5</sup>An implementation code in MATLAB of the Gradient Descent is available at <https://github.com/renvmorales/Linear-Regression>.

## 2.2 Feature Scaling

A typical problem that difficults the convergence of the Gradient Descent is the difference on the range of values of each feature. In fact, most variables used in prediction models can display a diversity in terms of unities and typical values. The computational gradient can then become dominated by a majority of small component values (lower scale feateures) and the size of each new “update step” can be severely decreased. A popular solution here consists of applying a “normalization” operation known as *feature scaling*. In feature scaling each corresponding feature  $\mathbf{x}_j$ , for  $j = 1, 2, \dots, n$ , is transformed into a new variable  $\widetilde{\mathbf{x}}_j$  according to

$$\widetilde{\mathbf{x}}_j = \frac{\mathbf{x}_j - \mu_{x_j}}{\sigma_{x_j}} \quad j = 1, 2, \dots, n \quad (9)$$

which is a rescaled version of the original feature with its mean value  $\mu_{x_j}$  and standard deviation  $\sigma_{x_j}$ . Ideally, using equation (9) will reduce the range of values to the  $[-1, 1]$  interval producing a zero mean feature<sup>6</sup>. Another common choice for the normalization factor is the largest deviation value  $\max\{\mathbf{x}_j\} - \min\{\mathbf{x}_j\}$  in the place of the standard deviation value.

A final observation is that using the Gradient Descent (or any other regression method) the transformation of variables will obviously produce a different optimized solution  $\widetilde{\theta}_0, \widetilde{\theta}_1, \widetilde{\theta}_2, \dots, \widetilde{\theta}_n$ , since features are now different from the originals. However, it is possible to use the following set of equations,

$$\begin{aligned} \theta_0 &= \widetilde{\theta}_0 - \sum_{j=1}^n \frac{\mu_{x_j} \widetilde{\theta}_j}{\sigma_{x_j}} \\ \theta_k &= \widetilde{\theta}_k / \sigma_{x_k} \quad k = 2, 3, \dots, n \end{aligned} \quad (10)$$

which recovers the expected solution related to the original features  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$ .

## 2.3 Normal Equation Method

An analytical solution to the Multivariate Linear Regression problem is the *Normal Equation* which can be obtained explicitly by either, setting the cost function derivatives with respect to  $\theta_j$ ’s to zero, or using Linear Algebra arguments. The following result provides the theoretical optimal value,

$$\boldsymbol{\theta}_{min} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad (11)$$

where, once again,  $\mathbf{X} \in \mathbb{R}^{m \times (n+1)}$ ,  $\mathbf{y} \in \mathbb{R}^m$  and  $\boldsymbol{\theta}_{min} \in \mathbb{R}^{n+1}$ . Although the apparent ease of applying Normal Equation to find the parameter values, it is important to observe that this method can also be disadvantageous in some situations. The major issue here is when the number of features  $n$  is sufficiently large so computing the inverse matrix term  $(\mathbf{X}^T \mathbf{X})^{-1}$  can be slow or even impractical. The table 1 describes the main differences when considering Gradient Descent and Normal Equation methods. A practical rule of thumb here is to consider the Gradient Descent method whenever the number of features  $n$  is something in the order of  $10^5$  or greater [Ng, 2017].

---

<sup>6</sup>This is expected when  $\mathbf{x}_j$  is normally distributed for example.

Table 1: Comparison between the Multivariate Linear Regression methods.

Gradient Descent	Normal Equation
Need to choose $\alpha$	No need to choose $\alpha$
Need many iterations	No need to iterate
Feature scaling is recommended	No need of feature scaling
$O(kn^2)$	$O(n^3)$ , need to compute $(\mathbf{X}^T \mathbf{X})^{-1}$
Works well when $n$ is large	Slow if $n$ is large

As a final remark, there are some specific situations which will cause  $\mathbf{X}^T \mathbf{X}$  to be non-invertible. In effect, this will hold true if the resulting  $\mathbf{X}^T \mathbf{X}$  matrix possess linear dependent features (columns), or if the number of training examples  $m$  is less or equal than the number of features  $n$ . In both scenarios, the resulting matrix is non-invertible and called *singular* or *degenerate*. The recommendation here is to discard some redundant features, or consider a regularization technique.

### 3 Polynomial Regression

A simple modification of the hypothesis function can give more flexibility when fitting a training dataset. The “trick” here is to include polynomial terms to the basic model using original data features information. For instance, suppose the case where  $n = 2$  features and the corresponding linear regression model:  $\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2$ . We want to fit our training data to a general quadratic form such as,

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1 x_2 + \theta_4 x_1^2 + \theta_5 x_2^2$$

where the two models differ by just 3 nonlinear terms. However, this quadratic polynomial model can also be viewed as a 5-dimensional linear model. This is possible by just feeding the 2-dimensional linear model with the 3 new following features,

$$x_3 = x_1 x_2$$

$$x_4 = x_1^2$$

$$x_5 = x_2^2$$

Concretely, by “artificially” including these new features on the original dataset, we provide a sufficient condition to fit a training data to the hypothesis  $\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4 + \theta_5 x_5$ . Therefore, solving a *polynomial regression* problem is equivalent to solve a “virtual”<sup>7</sup> and higher dimensional linear regression problem. In this sense, all previous methods of finding the optimum parameters for the minimization problem are valid (Gradient Descent and the Normal Equation).

A few observations remain on the use of polynomial models:

**1. Deciding the polynomial degree:** There is a no well defined criterion for defining the polynomial order that best fit the training data. A common practice is to fit different (polynomial) models, and use a *cross-validation* set to select the model with lower cost function-error (or equivalently lower RMSE) in order to make predictions.

<sup>7</sup>The word *virtual* here is used in the context of the features, treating them as if they all come from independent measurements, where in reality they are not!

**2. The parameter number “escalade”:** Suppose that, instead of using a parabolic equation, we have chosen a cubic polynomial model. The general form for this model would be,

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1 x_2 + \theta_4 x_1^2 + \theta_5 x_2^2 + \theta_6 x_1^2 x_2 + \theta_7 x_1 x_2^2 + \theta_8 x_1^3 + \theta_9 x_2^3$$

which indicates a substantial increase of the number of parameters. However, in many realistic applications the number of features would be far greater than 2. Because of the new terms will include combinations of all original features, the resulting total number of parameters will grow as  $\sim n^d$ , where  $d$  is the polynomial degree. Even if the required number of parameters is not exceptionally large, it may be greater than the number of training examples  $m$ . A reasonable solution then is to subjectively choose a few terms (the most “relevant” ones). Another option is to work with another regression method, such as *Neural Networks*, that can represent well complex behavior without requiring a large number of parameters.

**3. Feature scaling:** Since we are introducing nonlinear features to the training dataset, a *feature scaling* procedure is definitely recommended in order to assure an efficient convergence to the optimum parameter values when using the Gradient Descent method (this is not the case when using the Normal equation approach).

**4. Overfitting:** There is a natural tendency towards *overfitting* (when the model is too flexible and random errors are taken into account to the model) specially when using high order polynomials. This is more likely to happen when the number of parameters is greater than the number of training examples  $m$ . In this case, a *regularization* procedure can be helpful by “smoothing” the model response (lowering the variance).

**5. Other nonlinear terms:** In theory, some other nonlinear terms based on the following functions,  $\sqrt{x}$ ,  $\sqrt[3]{x}$ ,  $1/x^2$ ,  $\exp(x)$ ,  $\log(-x)$  could be also included to the model. However, this is not a very common practice, for a large number of features  $m$  we mostly do not possess strong evidences that any of these terms will be effective adjusting data (except for the case  $m = 2$  where we can visually check this).

## 4 Regularization

Regularization is meant to control and reduce the overfitting problem: the model have adjusted excessively well available data, but fails on generalizing and making predictions based on new data. The idea is to preserve the model complexity (for instance, a high polynomial degree) while “penalizing” the parameter values on the cost function  $J$  as follows,

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] \quad (12)$$

The  $\lambda$  is known as the *regularization parameter* and allows to modulate the “inflation effect” over the parameter values. Concretely, since overfitting is usually related to low training errors, minimizing (12) would drastically reduce the parameter values that has more effect on the cost. The optimization problem can be restated as,

$$\min_{\theta \in \mathbb{R}^{n+1}} \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] \quad (13)$$

The next sections will add the regularization parameter to the Gradient Descent and the Normal Equation algorithms.

## 4.1 Gradient Descent with Regularization

In order to account for the regularization on the Gradient Descent update rule, an additional term (scaled by the learning rate  $\alpha$ ) should be added to the cost function gradient as follows,

$$\boldsymbol{\theta}_{new} = \boldsymbol{\theta}_{old} - \alpha \left[ \nabla J(\boldsymbol{\theta}_{old}) + \frac{\lambda}{m} \tilde{\boldsymbol{\theta}} \right] \quad (14)$$

where,  $\lambda$  is the chosen regularization parameter,  $m$  is the number of training examples, and  $\tilde{\boldsymbol{\theta}}$  is the exactly same  $\Re^{n+1}$  vector  $\boldsymbol{\theta}_{old}$ , except from its first component which is equal to zero. The following set of updating rules for each component can be expressed as the following assignment algorithm operations,

$$\begin{aligned} \theta_0 &:= \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) \\ \theta_j &:= \theta_j - \alpha \left[ \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \right] \quad j \in \{1, 2, \dots, n\} \end{aligned} \quad (15)$$

Concretely, the second update rule can be also expressed (after some manipulation) as,

$$\theta_j := \theta_j \left( 1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)} \quad j \in \{1, 2, \dots, n\}$$

where the term  $1 - \alpha \frac{\lambda}{m}$  should be less than 1 and the second term stays the same as in the original update rule from equation (7). The regularization parameter has a more pronounced decreasing effect on the parameter values when compared to the non regularized case ( $\lambda = 0$ ).

The Algorithm 2 describes the Gradient Descent with the regularization procedure taken into account<sup>8</sup>

---

### Algorithm 2 Gradient Descent Algorithm with Regularization

---

- 1: Set initial guess for the parameter values:  $\theta_0, \theta_1, \theta_2, \dots, \theta_n$
  - 2: Set the learning rate value  $\alpha$
  - 3: Set the total number of update iterations  $N_{\text{iter}}$
  - 4: Set the regularization parameter  $\lambda$
  - 5: **for**  $k = 1, 2, \dots, N_{\text{iter}}$  **do**
  - 6:   Compute the parameters:
 
$$\begin{aligned} \theta_j &:= \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}, & j = 0 \\ \theta_j &:= \theta_j - \alpha \left[ \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \right], & j = 1, 2, \dots, n \end{aligned}$$
  - 7:   Compute the cost
 
$$J(k) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$
  - 8: **end for**
- 

<sup>8</sup>An implementation code in MATLAB of the Gradient Descent with regularization is available at <https://github.com/renvmorales/Linear-Regression>.

## 4.2 Normal Equation with Regularization

The non iterative Normal Equation procedure can add in the regularization parameter inside the parentheses of the original closed form,

$$\boldsymbol{\theta}_{min} = \left( \mathbf{X}^T \mathbf{X} + \lambda \mathbf{L} \right)^{-1} \mathbf{X}^T \mathbf{y} \quad (16)$$

where the matrix  $\mathbf{L} \in \Re^{(n+1) \times (n+1)}$  is very much alike the *identity matrix*, except by the fact the first diagonal element is zero,

$$\mathbf{L} = \begin{bmatrix} 0 & & & & \\ & 1 & & & \\ & & 1 & & \\ & & & \ddots & \\ & & & & 1 \end{bmatrix}$$

Computing the expression (16) with a positive regularization parameter can help to overcome the problem: *if  $m \leq n$ , then  $\mathbf{X}^T \mathbf{X}$  is not invertible*. In fact, even if that is the case, the term  $\mathbf{X}^T \mathbf{X} + \lambda \mathbf{L}$  can become invertible with the right choice of  $\lambda$ .

## References

- Joseph B. Keller. Inverse problems. *The American Mathematical Monthly*, 83:107–118, 1976. URL <http://www.maa.org/programs/maa-awards/writing-awards/inverse-problems>.
- Andreas Kirsch. *An introduction to the mathematical theory of inverse problems*, volume 120. Springer Science & Business Media, 2011.
- Andrew Ng. Machine learning online open course at coursera, 2017. URL <https://www.coursera.org/learn/machine-learning/>.